

```

<!-- TOC -->
* [General](#general)
* [Create pods](#create-pods)
* [Test a pod](#test-a-pod)
* [Pods, containers and storage](#pods-containers-and-storage)
* [Create other resources](#create-other-resources)
* [Update resources](#update-resources)
* [Debugging](#debugging)
* [Delete / replace resources](#delete--replace-resources)
* [Secrets for ServiceAccount](#secrets-for-serviceaccount)
* [Networking, services, DNS](#networking-services-dns)
* [Cluster](#cluster)
* [Helm](#helm)
* [etcd](#etcd)
<!-- TOC -->

### General
- In Mousepad, prepare these setup commands:
```
git clone https://github.com/glahitette/kubernetes
cd kubernetes/
chmod +x bashrc_append.sh
./bashrc_append.sh
```
- To assume root privileges: `sudo -i` (only applies to ACG server setup?)
- Use `grep -A2 Mounts` to show two lines after the line matching `Mounts`
- Select the `acgk8s` cluster to interact: `k config use-context acgk8s`
- Watch pods / deployments / jobs: `k get pods -w` / `k get deployments -w` / `k get jobs -w`
- Repeat command every n seconds, example: `watch -n 2 kubectl get pods`
- Check all resources [in all namespaces]: `k get all [-A]`
- List k8s "internal" pods, sorted by node name: `k get pods -n kube-system --sort-by .spec.nodeName`
- List of resources: `k api-resources`
- API e.g. for pod manifests : `k explain pods[child1.child2] | more`

### Create pods
- Create an nginx pod: `k run my-pod --image=nginx [--port=80] [--labels app=b]`
- Create a busybox pod: `k run my-pod --image=busybox $do --command -- sh -c "touch /tmp/ready && sleep 1d" > my-pod.yml`
  - Command YAML syntax example: `command: ['sh', '-c', 'while true; do echo success > /output/output.log; sleep 5; done']`
- Create a throw-away, interactive pod with busybox | netshoot: `k run my-pod --image=(busybox | nicolaka/netshoot) --restart=Never --rm -ti`

### Test a pod
- With a command: `k exec my-pod [-c my-container] (-- env ... | -- cat ...)`
- In interactive mode: `k exec my-pod [-c my-container] -ti -- sh`

### Pods, containers and storage
- Startup probes: run at container startup and stop running once they succeed; very similar to liveness probes (which run constantly on a schedule); useful for legacy applications that can have long startup times.
- Readiness probes: used to prevent user traffic from being sent to pods that are still in the process of starting up (e.g. pod STATUS = Running but READY = "0/1")
  - Example: for a service backed by multiple container endpoints, user traffic will not be sent to a particular pod until its containers have all passed readiness checks.
- Pod's restart policy: Always (by default), OnFailure (restarted only if error code returned), and Never.
- Pod with InitContainer(s) will show "Init(0/n)" in their STATUS during initialisation
- A Mirror Pod represents a Static Pod in the Kubernetes API, allowing you to easily view the Static Pod's status.
- Volumes:
  - A PersistentVolume's `persistentVolumeReclaimPolicy` determines how the storage resources can be reused when the PersistentVolume's associated PersistentVolumeClaims are deleted:
    - `Retain`: Keeps all data. This requires an administrator to manually clean up the data and prepare the storage resource for reuse.
    - `Delete`: Deletes the underlying storage resource automatically (cloud only).
    - `Recycle`: Automatically deletes all data in the underlying storage resource, allowing the PersistentVolume to be reused.
  - `allowVolumeExpansion` property of a **StorageClass**, if set to false (per default), prevents from resizing a PersistentVolumeClaim.

### Create other resources
- Create a job with `k create job my-job --image=busybox $do > job.yml -- sh -c "sleep 2 && echo done"` then check pod execution (no such thing as starting a Job or CronJob!)

```

- Create a ConfigMap from a file, with a specific key: ``k create configmap my-cm --from-file=index.html=/opt/course/15/web-moon.html``
- Create a secret (with implicit base64 encoding): ``k create secret generic my-secret --from-literal user=test --from-literal pass=pwd``
- Create an nginx deployment: ``k create deployment my-dep --image=nginx $do > my-dep.yml`` (deployment name is used as prefix for pods' name)
- Create a Service...
 - ...to expose a given pod ``k expose pod my-pod --name my-svc --port 3333 --target-port 80`` (faster than creating a service and editing its selector labels)
 - ...for an nginx deployment, which serves on port 80 and connects to the containers on port 8000: ``k expose deployment nginx --port=80 --target-port=8000 [--type ClusterIP|NodePort|...] [$do]``
 - Note: A NodePort Service kind of lies on top of a ClusterIP one, making the ClusterIP Service reachable on the Node IPs (internal and external).
- Create a quota: ``k create quota my-q --hard=cpu=1,memory=1G,pods=2,services=3... [$do]``
- Create Role / ClusterRole to permissions within a namespace / cluster-wide: ``k create role my-role --verb=get,list,watch --resource=pods,pods/logs``
- Create RoleBinding / ClusterRoleBinding to connect Roles / ClusterRoles to subjects (users, groups or ServiceAccounts): ``k create rolebinding my-rb --role=my-role --user=my-user``
- Create a service account to allow container processes within Pods to authenticate with the K8s API: ``k create sa my-sa``

Update resources

- Add / remove / change a label: ``k label pods my-pod app=b`` / ``k label pods my-pod app`` / ``k label pods my-pod app=v2 --overwrite``
- Add a new label tier=web to all pods having 'app=v2' or 'app=v1' labels: ``k label po -l "app in(v1,v2)" tier=web``
- Change a pod's image: ``k set image my-pod nginx=nginx:1.7.1 [--record]``
- Recreate the pods in a deployment: ``k rollout restart deploy my-dep``
- Perform a rolling update (e.g. to change an image): ``k edit deployment my-dep`` or ``k set image deployment my-dep nginx=nginx:1.21.5 [--record]``
- Check rollout status: ``k rollout status deployment my-dep``
- Roll back to the previous version: ``k rollout undo deployment my-dep``
- Scale a deployment [and record the command (into Annotations > change-cause)]: ``k scale deployment my-dep --replicas=5 [--record]``
- Autoscale a deployment, pods between 5 and 10, targetting CPU utilization at 80%: ``k autoscale deploy my-dep --min=5 --max=10 --cpu-percent=80``
- View the Horizontal Pod Autoscalers (hpa): ``k get hpa nginx``

Debugging

- Use ``k get pods [-A] [--show-labels]``: check ``STATUS``, ``READY`` and ``RESTARTS`` attributes.
- Retrieve a pod status: ``k get pod <pod_name> -o json | jq .status.phase``
- Retrieve pod / container logs: ``k logs <pod_name> [-c <container_name>] [-p]`` (if pod crashed and restarted, -p option gets logs about the previous instance)
- List events for a / all namespace(s): ``k get events (-n <my-namespace> | -A)``
- Show metrics for pod(s) / nodes: ``k top pod [--containers] [--sort-by (cpu | memory)] [-l app=b]`` / ``k top node [--sort-by (cpu | memory)]``

Delete / replace resources

- Force replace a resource: ``k replace --force -f ./pod.json``
- Delete pods and services using their label: ``k delete pods,services -l app=b $now``

Secrets for ServiceAccount

- If a Secret belongs to a ServiceAccount, it'll have the annotation ``kubernetes.io/service-account.name``
- Use ``k get secret ...`` to get a base64 encoded token
- Use ``k describe secret ...`` to get a base64 decoded token...or pipe it manually through ``echo <token> | base64 -d -``

Networking, services, DNS

- The cluster has a single virtual network spanning across all Nodes.
- Kubernetes `**nodes**` will remain ``NotReady``, unable to run Pods, until a network plugin is installed. ``Starting kube-proxy`` will be shown in the nodes logs and no networking pods will exist.
- Default FQDN:
 - ``<pod-ip-address-with-dashes>.<my-namespace>.pod.cluster.local.``
 - ``<my-service-name>.<my-namespace>.svc.cluster.local.``
- ``from`` and ``to`` selectors:
 - ``![(np_from_to_selectors.png)]``
- Example of "deny all" policy for labelled pods:


```

...

```

spec:

```

podSelector:
  matchLabels:

```

```

    app: maintenance
  policyTypes:
    - Ingress
    - Egress
  ...
- Example of egress policy, 1) restricting outgoing tcp connections from frontend to api
, 2) still allowing outgoing traffic on UDP/TCP ports 53 for DNS resolution.
...
spec:
  podSelector:
    matchLabels:
      id: frontend          # label of the pods this policy should be applied on
  policyTypes:
    - Egress                # we only want to control egress
  egress:
    - to:                   # 1st egress rule
      - podSelector:        # allow egress only to pods with api label
        matchLabels:
          id: api
    - ports:                # 2nd egress rule
      - port: 53             # allow DNS UDP
        protocol: UDP
      - port: 53             # allow DNS TCP
        protocol: TCP
  ...
- Example of Network Policy allowing all pods in the `users-backend` namespace to
communicate with each other only on a specific port (80): first label the namespace: `k
label namespace users-backend app=users-backend` then use:
...
metadata:
  name: np-users-backend-80
  namespace: users-backend
spec:
  podSelector: {} # selects all pods in the specified namespace
  policyTypes:
    - Ingress
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            app: users-backend
      ports:
        - protocol: TCP
          port: 80
  ...
- Example policy containing a single `from` element allowing connections from Pods with
the label `role=client` in namespaces with the label `user=alice`
...
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            user: alice
      podSelector:
          matchLabels:
            role: client
  ...
- Example policy containing two elements in the `from` array, allowing connections from
Pods in the local Namespace with the label `role=client`, **or** from any Pod in any
namespace with the label `user=alice`.
...
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            user: alice
      - podSelector:
          matchLabels:
            role: client
  ...
- When in doubt, `kubectl describe` shows how Kubernetes has interpreted the policy.
- Endpoints are the underlying entities (such as Pods) that a Service routes traffic to.
- Ingress: manages external access to Services; more powerful than a simple NodePort
Service (e.g. SSL termination, advanced load balancing, or namebased virtual hosting).

### Cluster
- Drain a node: `k drain [--ignore-daemonsets --force] <node name>`

```

- The `kubectl drain` subcommand on its own does not actually drain a node of its DaemonSet pods: the DaemonSet controller (part of the control plane) immediately replaces missing Pods with new equivalent Pods.
- The DaemonSet controller also creates Pods that ignore unschedulable taints, which allows the new Pods to launch onto a node that you are draining.
- Resume scheduling `**new pods**` onto the node: `k uncordon <node name>`
- In a cluster built with `kubeadm`:
 - Check the status of cluster components (e.g. kube-apiserver): check the status of (static) Pods in the `kube-system` Namespace (kube-apiserver is not set up as a systemd service).
 - Find logs for the Kubernetes API Server: `k logs -n kube-system <api-server-pod-name>` (the `/var/log/kube-apiserver.log` log file is not available on the host since the API Server runs in a static Pod).
 - Find kubelet logs: `sudo journalctl -fu kubelet` (kubelet runs as a standard service).
 - Investigate DNS issues: check the DNS Pods in the `kube-system` Namespace.
- Upgrade `kubeadm` clusters: [\[link\]\(CKA%20training/Upgrading%20kubeadm%20clusters.md\)](#)

Helm

- List release with `helm [-n my_ns] ls [-a]`
- List pending deployments on all namespaces: `helm list --pending -A`
- List / search repo: `helm repo list` / `helm search repo nginx`
- Download (not install) a chart from a repository: `helm pull [chart URL | repo/chartname] [...] [flags]`
- Untar a chart (after downloading it): `helm pull --untar [rep/chartname]`
- Check customisable values setting for an install, e.g. `helm show values bitnami/apache [yq e]`
- Custom install example `helm install my-apache bitnami/apache --set replicaCount=2`
- Upgrade a release, e.g. `helm upgrade my-api-v2 bitnami/nginx`
- Undo a helm rollout/upgrade: `helm rollback`
- Delete an installed release with `helm uninstall <release_name>`

etcd

- etcd is a consistent and highly-available key value store used to store for all cluster data.
- Backup / restore etcd data: [\[link\]\(CKA%20training/etcd.md\)](#)